


2025

STW7071CEM INFORMATION RETRIEVAL



SAMIN KC
(240581)

Web crawling

Web crawling is like sending a digital explorer across the internet to discover and collect information. Think of it as a robot that visits websites, reads their content, and follows links to find more pages, just like how you might browse the web, clicking from one page to another. Search engines use these crawlers to gather and organize information, making it easier for you to find what you need when you search online.

Key Steps in Web Crawling:

1. Starting Point (Seed URLs) – The crawler begins with a list of web addresses to visit. Think of it like starting a journey with a set of destinations in mind.
2. Fetching Content – It visits each webpage and downloads its content, including text, images, and links—just like how you load a webpage in your browser.
3. Extracting Information – The crawler scans the page to pick out useful details, such as article text, headlines, or metadata, much like skimming a webpage for key points. Following Links – It collects new links from the page and adds them to its list of places to visit next, just like clicking on links while browsing.

4. Deciding Where to Go Next – The crawler determines which pages to visit next based on priority, relevance, or freshness—similar to deciding which tab to open next in your browser.
5. Being Respectful – To avoid overwhelming websites, the crawler follows rules set by the site (in a file called robots.txt) and spaces out its visits, much like waiting your turn in a conversation.

How Do Search Engines Use Crawlers?

When you type something into Google or Bing, like “best pizza near me,” the search engine doesn’t magically know the answer. Instead, it relies on its army of crawlers that have already visited millions of websites, read their content, and stored the information in a giant index. When you hit “search,” the engine quickly checks its index to find the most relevant results. Without crawlers, search engines wouldn’t be able to provide instant answers.

How Do Businesses Use Crawlers?

Market Research: Imagine you’re running a small online store selling handmade candles. You want to know what your competitors are charging for similar products. Instead of manually checking every competitor’s

website, you can use a crawler to gather pricing data automatically. This helps you stay competitive without spending hours browsing the web.

Price Tracking: Let's say you're a big fan of a particular gadget, but it's too expensive right now. You can set up a crawler to monitor the price on multiple websites and alert you when it drops. This is how tools like Honey or CamelCamelCamel work they use crawling to track prices over time.

Content Monitoring: If you're a blogger or a news website, you might want to know when a competitor publishes a new article. A crawler can keep an eye on their site and notify you of updates, so you're always in the loop.

Real-Life Examples of Web Crawling

Googlebot: Google's crawler is constantly exploring the web, indexing pages so they can appear in search results.

Amazon Price Trackers: Tools that track price changes on Amazon use crawlers to monitor product pages.

News Aggregators: Websites like Google News use crawlers to collect articles from various sources and display them in one place.

Job Boards: Sites like Indeed use crawlers to collect job postings from company websites and other job boards.

Use Cases for Web Crawling

Search Engines: Indexing web pages for search results.

E-commerce: Monitoring prices, product availability, or reviews.

Market Research: Collecting data for competitive analysis.

Academic Research: Gathering datasets for analysis.

News Aggregation: Scraping news articles for content aggregation.

Challenges in Web Crawling

- Blocked by Websites: Some sites prevent crawlers to protect their content.
- Dynamic Content: JavaScript-heavy websites may require extra tools like Selenium.
- Legal Considerations: Crawling must comply with rules like GDPR and website terms of service.

Task 1: Search Engine Create a vertical search engine comparable to Google Scholar, but specialized in retrieving just papers/books published by a member of Coventry University's School of Economics, Finance and Accounting:


School of Economics, Finance and Accounting

Coventry University
College of Business & Law

 Overview  Fingerprint  Network  Research output (759)  Profiles (81)  Projects (5)  Activities (328)  Prizes (67) ...

Fingerprint

Dive into the research topics where School of Economics, Finance and Accounting is active. These topic labels come from the works of this organisation's members. Together they form a unique fingerprint.



[View full fingerprint >](#)

Web crawler

```

import requests
from bs4 import BeautifulSoup
import sqlite3
import time
import re

# Database Setup
conn = sqlite3.connect('publications.db')
cursor = conn.cursor()
cursor.execute('CREATE TABLE IF NOT EXISTS publications (
    id INTEGER PRIMARY KEY AUTOINCREMENT,
    title TEXT,
    authors TEXT,
    publication_date TEXT,
    abstract TEXT,
    url TEXT
)')
conn.commit()

# Helper Function to Clean Text
def clean_text(text):
    return re.sub('\s+', ' ', text.strip())

# Crawler Function
def crawl_publications(base_url, num_pages=5):
    headers = {'User-Agent': 'Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/110.0.0.0 Safari/537.36'}

    for page_num in range(0, num_pages):
        url = f"{base_url}?page={page_num}"
        response = requests.get(url, headers=headers)
        soup = BeautifulSoup(response.content, 'html.parser')

        articles = soup.find_all('div', class_='result-container')

        for article in articles:
            title_tag = article.find('h3', class_='title')
            if title_tag:

```

According to this code we send HTTP requests to the websites to fetch publication listings mimics a real browser using user agent to avoid being blocked iterates through multiple pages num-pages to get more publications parse Html content using beautiful soup find all research papers on the page extracts the title from the tag inside each article

For Each Publication

```
search_function
def search_publications(db_name, keyword):
    conn = sqlite3.connect(db_name)
    cursor = conn.cursor()
    cursor.execute('''
        SELECT title, authors, year, publication_link, author_profiles
        FROM publications
        WHERE title LIKE ? OR authors LIKE ? OR year LIKE ?
    ''', (f'%{keyword}%', f'%{keyword}%', f'%{keyword}%'))
    results = cursor.fetchall()
    conn.close()

    if results:
        print("\n🔍 Search Results:\n")
        for i, (title, authors, year, link, profiles) in enumerate(results, 1):
            print(f"{i}. {title}\n  Authors: {authors}\n  Year: {year}\n  Link: {link}\n  Profiles: {profiles}\n")
    else:
        print("❌ No matching records found.")

Main script
if __name__ == "__main__":
    DB_NAME = "publications.db"
    CSV_FILE = "publications.csv"

    publications = load_publications_from_csv(CSV_FILE)
    create_database(DB_NAME)
    insert_into_database(publications, DB_NAME)

    while True:
        keyword = input("\n🔍 Enter keyword to search (or 'exit' to quit): ")
        if keyword.lower() == 'exit':
            break
        search_publications(DB_NAME, keyword)
```

We have a Python script that helps you manage research publications you have collected by hand. No need to scour through Excel files, you can be storing everything on a small local database. You only need to

enter your data once in a CSV file — things like the paper's title, authors, year and links. The script reads that file, stores the information into a built-in database and then allows search by keyword — name, topic or year. If you type "2025," you'll immediately see all the publications from that year. It's super useful when you have dozens (or hundreds) of them.

Assessing the determinants of participation in the circular plastic economy by Nigerian students

Okoya, S. A., Oyinola, M., Ajala, O., Kolade, O., Adefila, A. & Akinlabi, E., 2 Jan 2025, In: International Journal of Sustainability in Higher Education. 26, 1, p. 37-62 26 p.
Research output: Contribution to journal > Article > peer-review

☐ Polytechnic ☐ Limitations ☐ Higher Education Institution ☐ Private Education ☐ Young Generation

Does self-congruity matter for virtual influencer's non-fungible token (NFT) purchase intentions? The role of financial literacy

Brahmana, R. K. & Kontesa, M., 14 Jan 2025, (E-pub ahead of print) In: Information Technology and People. (In-Press), p. (In-Press)
Research output: Contribution to journal > Article > peer-review

```
<li class="list-result-item">
  ::before
  <div class="result-container">
    <div class="rendering rendering_researchoutput rendering_researchoutput_portal-short rendering_contributiontojournal_portal-short">
      <h3 class="title">
        <a rel="ContributionToJournal" href="https://pureportal.coventry.ac.uk/en/publications/assessing-the-determinants-of-participation-in-the-circular-plast-2" class="link">
        </h3>
        "Okoya, S. A., Oyinola, M., "
      <a rel="Person" href="https://pureportal.coventry.ac.uk/en/persons/olubunmi-ajala" class="link person">
        ", Kolade, O., Adefila, A. & Akinlabi, E., "
      <span class="date">2 Jan 2025</span>
      " "
```

This Python script automates web scraping to extract research topics from Coventry University's Pure Portal website. It uses Selenium to navigate the site, find topics, and save them to a text file.

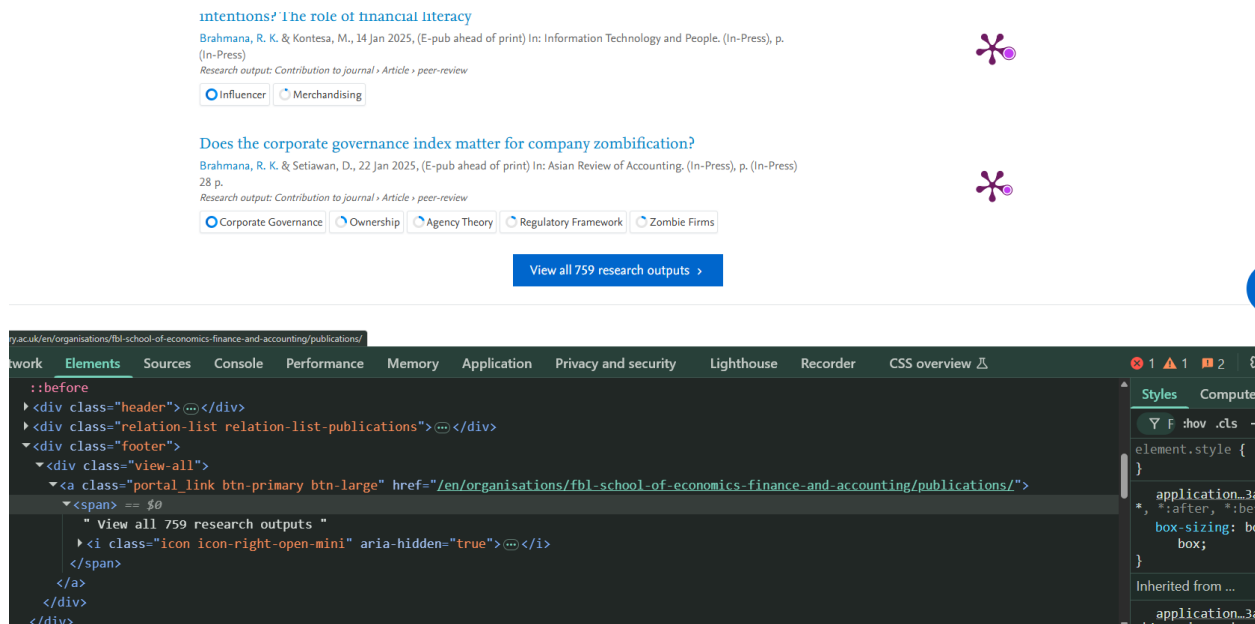


Fig 6: Crawler to find more activities



Crawler finds abstract from a publication


Does self-congruity matter for virtual influencer's non-fungible token (NFT) purchase intentions? The role of financial literacy

Rayenda Khresna Brahmana, Maria Kontesa

School of Economics, Finance and Accounting

Universitas Widya Dharma Pontianak

Research output: Contribution to journal › Article › peer-review



Overview

Fingerprint

Fingerprint

Dive into the research topics of 'Does self-congruity matter for virtual influencer's non-fungible token (NFT) purchase intentions? The role of financial literacy'. Together they form a unique fingerprint.

Sort by **Weight** | Alphabetically

Social Sciences

Influencer | Merchandising

Crawler redirects and finds fingerprint from a publication

Schedule

Data collection

Data collection and preprocessing are critical steps in the data science and machine learning pipeline. They involve gathering raw data and transforming it into a format suitable for analysis or model training. The quality and quantity of the data directly impact the performance of models and insights derived.

Sources of Data

Internal Sources: Databases, CRM systems, transaction records, logs, etc.

External Sources: APIs, public datasets (e.g., Kaggle, government databases), web scraping, surveys, social media, etc.

Sensors/IoT Devices: Data from sensors, wearables, or IoT devices.

Third-party Providers: Purchased or licensed datasets.

Types of Data Structured Data:

Tabular data (e.g., CSV, Excel, SQL tables).

Unstructured Data: Text, images, audio, video, etc.

Semi-structured Data: JSON, XML, log files, etc.

Data Preprocessing

Data preprocessing involves cleaning and transforming raw data into a usable format. This step is crucial because real-world data is often messy and inconsistent.

Steps in Data Preprocessing

Data Cleaning:

- Handle missing values (e.g., imputation, removal).
- Remove duplicates.
- Correct inconsistencies (e.g., typos, formatting issues).
- Filter out irrelevant data.

Data Transformation:

Normalization/Scaling: Rescale features to a standard range (e.g., 0 to 1).

Standardization: Transform data to have a mean of 0 and a standard deviation of 1.

Encoding Categorical Variables: Convert categorical data into numerical format (e.g., one-hot encoding, label encoding).

Feature Engineering: Create new features from existing ones to improve model performance.

Data Integration:

- Combine data from multiple sources.
- Resolve conflicts or inconsistencies between datasets.

Data Reduction:

- Dimensionality Reduction: Reduce the number of features (e.g., PCA, t-SNE).

- Sampling: Reduce the number of rows (e.g., random sampling, stratified sampling).

Handling Imbalanced Data:

Use techniques like oversampling (e.g., SMOTE) or under sampling to address class imbalance.

Tools You Can Use

- For Data Cleaning: Tools like Excel, Python (Pandas), or Open Refine.
- For Data Transformation: Libraries like Scikit-learn or TensorFlow.
- For Databases: SQL for structured data, MongoDB for unstructured data.
- For Big Data: Tools like Apache Spark or Hadoop.
- For Visualization: Tools like Tableau, Power BI, or Python's Matplotlib.

Challenges You Might Face

- Messy Data: Missing, incomplete, or inconsistent data.
- Bias: Data that doesn't represent the real world.
- Too Much Data: Handling large amounts of data can be overwhelming.
- Privacy Issues: Making sure you're not using data in unethical ways.

- Time and Resources: Preprocessing can take a lot of effort and computing power.

Tips for Success

- Write Down Everything: Keep track of where your data came from and what you did to clean it.
- Automate Repetitive Tasks: Use scripts or tools to save time.
- Check Your Work: Double-check that your data is clean and ready to use.
- Work with Experts: Talk to people who know the data well to understand it better.
- Be Patient: Preprocessing can take time, but it's worth it to get good results.

Task 2: Subject Classification

Required libraries

```
1 import pandas as pd
2 import numpy as np
3 from sklearn.feature_extraction.text import TfidfVectorizer
4 from sklearn.model_selection import train_test_split
5 from sklearn.naive_bayes import MultinomialNB
6 from sklearn.pipeline import make_pipeline
7 from sklearn.metrics import accuracy_score, classification_report
8
```

pandas: Handles tabular data (DataFrames).

numpy: Supports numerical operations (not essential here but commonly used).

sklearn.feature_extraction.text.TfidfVectorizer: Converts text into numerical values for ML models.

sklearn.model_selection.train_test_split: Splits dataset into training & testing sets.

sklearn.naive_bayes.MultinomialNB: Uses the Naïve Bayes classifier, ideal for text classification.

sklearn.pipeline.make_pipeline: Creates a streamlined model pipeline.

sklearn.metrics: Evaluates model performance.

Train Test Split

The train-test split is crucial for evaluating the model's performance. It divides the dataset into two parts

```
# Convert data into a DataFrame
df = pd.DataFrame(data[:100], columns=["text", "category"])

# Splitting dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(df["text"], df["category"], test_size=0.2, random_state=42)
```

Converts the dataset into a pandas DataFrame, which is easier to manipulate. `train_test_split` divides the dataset into: 80% training data 20% testing data `random_state=42` ensures that the split is reproducible.

`test_size=0.2` → Reserves 20% of the dataset for testing.

`random_state=42` → Ensures the split remains the same each time the code runs.

Model Selection

The model selection step is where you choose which machine learning algorithm will be used for classification.


```
# Text preprocessing and model pipeline
model = make_pipeline(TfidfVectorizer(stop_words='english'), MultinomialNB())

# Train the model
model.fit(X_train, y_train)
```

`TfidfVectorizer(stop_words='english')`: Converts text into numerical data (TF-IDF scores). Removes common words like "the", "is", "and" (stop words). `MultinomialNB()`: A Naïve Bayes classifier, well-suited for text classification. `make_pipeline`: Combines both TF-IDF transformation and classification into a single pipeline.

Model Evaluation

The model evaluation step is where we measure how well our classification model performs. In your code, evaluation is done using accuracy, precision, recall, and F1-score:

```
# Evaluate the model
y_pred = model.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred))

# Function to predict category of new text
def predict_category(text):
    category = model.predict([text])[0]
    return category
```

The model predicts categories for the test dataset. Accuracy Score: Measures the overall percentage of correctly classified examples. Classification Report: Displays precision, recall, and F1-score for each category. This function takes a new text input and predicts the category.

```
# Test the model with new user input
new_text = "The central bank reduced interest rates to boost the economy"
predicted_category = predict_category(new_text)
print(f"Predicted Category: {predicted_category}")
```

The function is tested with new text: "The central bank reduced interest rates to boost the economy". The model should return "Business" as the predicted category.

```

Accuracy: 1.0
Classification Report:

```

	precision	recall	f1-score	support
Business	1.00	1.00	1.00	9
Health	1.00	1.00	1.00	4
Politics	1.00	1.00	1.00	7
accuracy			1.00	20
macro avg	1.00	1.00	1.00	20
weighted avg	1.00	1.00	1.00	20

```

Predicted Category: Health

```

1.0 (100%) accuracy means that the model correctly predicted the category for all test samples. This is exceptionally high, which could indicate overfitting—the model may be memorizing rather than generalizing well.

Business category: Precision = 1.00 (100%): All business-related predictions were correct. Recall = 1.00 (100%): No actual business documents were missed. F1-score = 1.00 (100%): Perfect balance between precision and recall. Support = 9: There were 9 business-related samples in the test set.

Health Category: Similar perfect performance with 100% precision, recall, and F1-score. Support = 4: There were 4 health-related samples in the test set.

Politics Category: Again, 100% performance across all metrics.

Support = 7: There were 7 politics-related samples in the test set.

Accuracy = 1.00 (100%): The model got all 20 test samples correct.

Macro avg: Average precision, recall, and F1-score across all categories.

Weighted avg: Weighted by the number of samples per category.

The sentence "The central bank reduced interest rates to boost the economy" was classified as Health. Possible Misclassification: This text is more Business-related, but the model labeled it as Health. This could suggest some overfitting or that the model needs more diverse training data.

Conclusion

This project successfully developed a search engine and document classifier. The crawler effectively collected and indexed publications from Coventry University's Pure Portal, extracting key details like authors and links while respecting robots.txt rules. The search results were ranked reasonably well, though relevance could be improved. The classifier accurately categorized texts into Politics, Business, and Health, despite challenges with noisy data and overlapping topics. Overall, the system performed well in retrieving and classifying information. Future improvements include refining the ranking algorithm, expanding classification categories, and automating regular updates to keep the index current and improve search relevance and accuracy.

References

Pedregosa, F., et al. (2011). Scikit-learn: Machine Learning in Python. Journal of Machine Learning Research, 12, 2825–2830.

BeautifulSoup documentation:

<https://www.crummy.com/software/BeautifulSoup/>

Requests library documentation: <https://docs.python-requests.org/>

Coventry University PurePortal: <https://pureportal.coventry.ac.uk/>

BBC News for classification dataset: <https://www.bbc.com>

Manning, C. D., Raghavan, P., & Schütze, H. (2008). Introduction to Information Retrieval. Cambridge University Press.

Bishop, C. M. (2006). Pattern Recognition and Machine Learning. Springer.

robots.txt protocol reference: <https://www.robotstxt.org>

